

Integration Testing Framework

This document describes the integration testing framework for the agent-orchestration-ops repository.

Overview

The integration testing framework provides comprehensive validation of the routing infrastructure, including:

- **Router health and basic functionality** (smoke tests)
- **Provider fallback behavior** (resilience tests)
- **Cost tracking and attribution** (ledger tests)

All tests run against a local Docker Compose environment that mirrors production infrastructure.

Architecture

Test Environment

The test environment is defined in `tests/compose.int.yml` and includes:

- **Redis** (7-alpine): Caching and state management
- **LiteLLM**: Router service with provider fallback logic
- **vLLM**: Local LLM inference engine (GPU-enabled)

All services include health checks and proper dependency ordering.

Test Suites

1. Smoke Tests (`test_router_smoke.py`)

Basic functionality validation:

- Health endpoint availability
- Simple chat completion requests
- Response format validation
- Model listing
- Metadata handling

Key function: `wait_health()` - Waits for router to become healthy before running tests.

2. Fallback Tests (`test_fallbacks.py`)

Provider resilience validation:

- Timeout handling and fallback routing
- Rate limit handling
- Fallback chain exhaustion
- Metadata preservation through fallbacks
- Latency impact of fallback routing

3. Cost Ledger Tests (`test_cost_ledgers.py`)

Cost tracking validation:

- Usage metrics structure (prompt_tokens, completion_tokens, total_tokens)
- Multi-tenant cost attribution

- Cost accumulation across requests
- Metadata preservation for attribution
- Streaming response cost tracking
- Edge cases (minimal tokens, consistency)

Running Tests Locally

Prerequisites

- Docker and Docker Compose installed
- Python 3.11+ with pip
- (Optional) NVIDIA GPU with Docker runtime for vLLM

Quick Start

1. Install test dependencies:

```
bash
pip install pytest requests
```

2. Start test environment:

```
bash
cd tests
docker compose -f compose.int.yml up -d
```

3. Wait for services to be ready (30-60 seconds):

```
bash
docker compose -f compose.int.yml ps
docker compose -f compose.int.yml logs -f litellm
```

4. Run all tests:

```
bash
pytest tests/ -v
```

5. Run specific test suite:

```
bash
pytest tests/test_router_smoke.py -v
pytest tests/test_fallbacks.py -v
pytest tests/test_cost_ledgers.py -v
```

6. Cleanup:

```
bash
docker compose -f tests/compose.int.yml down -v
```

Environment Variables

- `ROUTER_URL` : Router base URL (default: `http://localhost:4000`)
- `HEALTH_TIMEOUT` : Seconds to wait for health (default: `60`)

Example:

```
ROUTER_URL=http://localhost:4000 HEALTH_TIMEOUT=120 pytest tests/ -v
```

CI/CD Integration

GitHub Actions Workflow

The integration tests run automatically on:

- Push to `main`, `develop`, or `feature/**` branches
- Pull requests to `main` or `develop`

Workflow file: `.github/workflows/integration.yml`

Workflow Steps

1. **Checkout code**
2. **Set up Docker Buildx**
3. **Set up Python 3.11**
4. **Install pytest and requests**
5. **Start test environment** (docker compose up)
6. **Wait for services** (30s + health checks)
7. **Run smoke tests**
8. **Run fallback tests**
9. **Run cost ledger tests**
10. **Capture logs on failure** (redis, litellm, vllm)
11. **Upload logs as artifacts** (7-day retention)
12. **Cleanup** (docker compose down)

Viewing Test Results

- **GitHub Actions UI:** Check the “Actions” tab in the repository
- **Pull Request checks:** Tests must pass before merge
- **Logs on failure:** Download artifacts from failed runs

Test Development Guidelines

Adding New Tests

1. **Choose appropriate test file:**
 - `test_router_smoke.py` : Basic functionality
 - `test_fallbacks.py` : Resilience and error handling
 - `test_cost_ledgers.py` : Cost tracking and attribution

2. **Use the `wait_health()` fixture:**

```
python
@pytest.fixture(scope="module", autouse=True)
def ensure_router_ready():
    wait_health()
```

3. **Include descriptive docstrings:**

```
```python
def test_new_feature():
 """
 Test description here.
```

Verifies that:

1. First behavior
2. Second behavior

```
"""
...

```

#### 4. Use proper assertions:

```
python
assert response.status_code == 200, f"Request failed: {response.text}"
```

#### 5. Clean up resources:

- Tests should be idempotent
- No persistent state between tests
- Use unique trace\_ids for tracking

## Best Practices

- **Timeouts:** Always set reasonable timeouts on requests
- **Retries:** Use `wait_health()` pattern for flaky operations
- **Metadata:** Include tenant, trace\_id, user\_id for tracking
- **Assertions:** Be specific about what you're testing
- **Error messages:** Include context in assertion messages
- **Isolation:** Tests should not depend on each other

## Troubleshooting

---

### Common Issues

#### Services not starting

**Symptom:** Tests fail with connection errors

**Solution:**

```
Check service status
docker compose -f tests/compose.int.yml ps

View logs
docker compose -f tests/compose.int.yml logs litellm
docker compose -f tests/compose.int.yml logs redis

Restart services
docker compose -f tests/compose.int.yml restart
```

#### Health check timeouts

**Symptom:** `TimeoutError: Router not healthy after 60s`

**Solution:**

- Increase `HEALTH_TIMEOUT` environment variable
- Check if services are actually running
- Verify network connectivity between containers

#### GPU not available for vLLM

**Symptom:** vLLM fails to start or falls back to CPU

**Solution:**

- Ensure NVIDIA Docker runtime is installed
- Check GPU availability: `nvidia-smi`
- Modify `compose.int.yml` to remove GPU requirement for testing

**Port conflicts**

**Symptom:** Error: port already in use

**Solution:**

```
Find process using port
lsof -i :4000
lsof -i :6379
lsof -i :8000

Kill process or change ports in compose.int.yml
```

**Debug Mode**

Run tests with verbose output:

```
pytest tests/ -v -s --tb=long
```

View real-time logs:

```
docker compose -f tests/compose.int.yml logs -f
```

## Maintenance

---

### Updating Test Environment

When infrastructure changes:

1. Update `tests/compose.int.yml` to match production config
2. Update test assertions if API contracts change
3. Add new tests for new features
4. Update this documentation

### Performance Considerations

- Tests run in ~5-10 minutes in CI
- Local runs may be faster (cached images)
- GPU tests require appropriate hardware
- Consider test parallelization for large suites

### Future Enhancements

---

Planned improvements:

- [ ] Load testing with k6 or locust
- [ ] Performance regression detection
- [ ] Multi-region failover tests

- [ ] Chaos engineering scenarios
- [ ] Integration with monitoring/alerting
- [ ] Test data generation and fixtures
- [ ] Parallel test execution
- [ ] Test coverage reporting

## Support

---

For issues or questions:

1. Check this documentation
2. Review test logs and error messages
3. Check GitHub Issues for similar problems
4. Create a new issue with:
  - Test output
  - Service logs
  - Environment details
  - Steps to reproduce

---

**Last Updated:** 2025-09-30

**Maintainer:** Empire325Marketing DevOps Team